

```

#Jeffery Dirden
#W214801986
#ITAI-1371

import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LassoCV
from sklearn.svm import SVR
from sklearn.metrics import mean_squared_error, r2_score

# Generate a synthetic dataset
np.random.seed(42) # For reproducibility

# Assume 1000 samples with 10 features
num_samples = 1000
num_features = 10

# Create random features and a target variable
X = np.random.rand(num_samples, num_features) * 100 # Features scaled between 0 and 100
y = X[:, 0] * 0.5 + X[:, 1] * 0.3 + X[:, 2] * 0.2 + np.random.randn(num_samples) * 5 # Linear relationship with noise

# Convert to a DataFrame for compatibility with pandas
columns = [f'feature_{i}' for i in range(1, num_features + 1)]
data = pd.DataFrame(X, columns=columns)
data['margin'] = y # Add target variable

print("Synthetic Dataset Preview:")
print(data.head())

# Separate features and target
X = data.drop(['margin'], axis=1) # Drop target column
y = data['margin'] # Isolate the target variable

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Scale numerical features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Lasso Regression
print("\n--- Lasso Regression ---")
lasso = LassoCV(cv=5, random_state=0)
lasso.fit(X_train_scaled, y_train)

# Lasso Evaluation
y_pred_lasso = lasso.predict(X_test_scaled)
lasso_rmse = np.sqrt(mean_squared_error(y_test, y_pred_lasso))
lasso_r2 = r2_score(y_test, y_pred_lasso)
print(f"Lasso Regression RMSE: {lasso_rmse}")
print(f"Lasso Regression R^2: {lasso_r2}")

# Support Vector Machine Regression
print("\n--- SVM Regression ---")
svm = SVR(kernel='rbf')

# GridSearch for best parameters
param_grid = {'C': [0.1, 1, 10], 'gamma': [1, 0.1, 0.01]}
grid_search = GridSearchCV(svm, param_grid, cv=5, scoring='neg_mean_squared_error')
grid_search.fit(X_train_scaled, y_train)

# Best model from GridSearch
best_svm = grid_search.best_estimator_
print("Best SVM Parameters:", grid_search.best_params_)

# SVM Evaluation
y_pred_svm = best_svm.predict(X_test_scaled)
svm_rmse = np.sqrt(mean_squared_error(y_test, y_pred_svm))
svm_r2 = r2_score(y_test, y_pred_svm)
print(f"SVM Regression RMSE: {svm_rmse}")
print(f"SVM Regression R^2: {svm_r2}")

```

```
# Final Summary
```

```
print("\n--- Model Performance Summary ---")
```

```
print(f"Lasso RMSE: {lasso_rmse}, R^2: {lasso_r2}")
```

```
print(f"SVM RMSE: {svm_rmse}, R^2: {svm_r2}")
```

```
↳ Synthetic Dataset Preview:
```

	feature_1	feature_2	feature_3	feature_4	feature_5	feature_6	\
0	37.454012	95.071431	73.199394	59.865848	15.601864	15.599452	
1	2.058449	96.990985	83.244264	21.233911	18.182497	18.340451	
2	61.185289	13.949386	29.214465	36.636184	45.606998	78.517596	
3	60.754485	17.052412	6.505159	94.888554	96.563203	80.839735	
4	12.203823	49.517691	3.438852	90.932040	25.877998	66.252228	

	feature_7	feature_8	feature_9	feature_10	margin
0	5.808361	86.617615	60.111501	70.807258	54.449005
1	30.424224	52.475643	43.194502	29.122914	41.149443
2	19.967378	51.423444	59.241457	4.645041	42.564448
3	30.461377	9.767211	68.423303	44.015249	30.924632
4	31.171108	52.006802	54.671028	18.485446	27.208161

```
--- Lasso Regression ---
```

```
Lasso Regression RMSE: 5.304533013175062
```

```
Lasso Regression R^2: 0.8960878604006127
```

```
--- SVM Regression ---
```

```
Best SVM Parameters: {'C': 10, 'gamma': 0.01}
```

```
SVM Regression RMSE: 5.2842766850530625
```

```
SVM Regression R^2: 0.8968799600855704
```

```
--- Model Performance Summary ---
```

```
Lasso RMSE: 5.304533013175062, R^2: 0.8960878604006127
```

```
SVM RMSE: 5.2842766850530625, R^2: 0.8968799600855704
```